

Written By: Isis Vasquez

4/3/2025

## Imports

```
# Basic libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# AutoViz
!pip install autoviz
from autoviz.AutoViz_Class import AutoViz_Class

# PyCaret
!pip install pycaret
from pycaret.classification import *

# Warnings
import warnings
warnings.filterwarnings('ignore')
```



Show hidden output

## Input Data

```
try:
    df = pd.read_csv('https://raw.githubusercontent.com/Ivasquez2003/Data-Mining-Project/re
    display(df.head())
except FileNotFoundError:
    print("Error: 'test_Y3wMUE5_7gLdaTN.csv' not found. Please upload the file or provide t
except pd.errors.EmptyDataError:
    print("Error: 'test_Y3wMUE5_7gLdaTN.csv' is empty.")
except pd.errors.ParserError:
    print("Error: Could not parse 'test_Y3wMUE5_7gLdaTN.csv'. Please check the file format.
except Exception as e:
    print(f"An unexpected error occurred: {e}")
```



	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome
0	10001002	Male	No	0	Graduate	No	5810

0	LP001002	Male	No	0	Graduate	No	3049
1	LP001003	Male	Yes	1	Graduate	No	4583
2	LP001005	Male	Yes	0	Graduate	Yes	3000
3	LP001006	Male	Yes	0	Not Graduate	No	2583
4	LP001008	Male	No	0	Graduate	No	6000

## ✎ Exploring Data

```
# Statistical Summary
df.describe()
```



	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_Histo
<b>count</b>	614.000000	614.000000	592.000000	600.000000	564.0000
<b>mean</b>	5403.459283	1621.245798	146.412162	342.000000	0.8421
<b>std</b>	6109.041673	2926.248369	85.587325	65.12041	0.3648
<b>min</b>	150.000000	0.000000	9.000000	12.000000	0.0000
<b>25%</b>	2877.500000	0.000000	100.000000	360.000000	1.0000
<b>50%</b>	3812.500000	1188.500000	128.000000	360.000000	1.0000
<b>75%</b>	5795.000000	2297.250000	168.000000	360.000000	1.0000
<b>max</b>	81000.000000	41667.000000	700.000000	480.000000	1.0000

```
# Summary Statistics
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null   object
1   Gender                601 non-null   object
2   Married               611 non-null   object
3   Dependents            599 non-null   object
4   Education             614 non-null   object
5   Self_Employed         582 non-null   object
```

```

5  Loan_Employed      614 non-null    object
6  ApplicantIncome    614 non-null    int64
7  CoapplicantIncome  614 non-null    float64
8  LoanAmount         592 non-null    float64
9  Loan_Amount_Term   600 non-null    float64
10 Credit_History     564 non-null    float64
11 Property_Area      614 non-null    object
12 Loan_Status        614 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB

```

```

# Analysis of Nulls
df.isnull().sum()

```

```


```

	0
<b>Loan_ID</b>	0
<b>Gender</b>	13
<b>Married</b>	3
<b>Dependents</b>	15
<b>Education</b>	0
<b>Self_Employed</b>	32
<b>ApplicantIncome</b>	0
<b>CoapplicantIncome</b>	0
<b>LoanAmount</b>	22
<b>Loan_Amount_Term</b>	14
<b>Credit_History</b>	50
<b>Property_Area</b>	0
<b>Loan_Status</b>	0

```
dtype: int64
```

```

# AutoViz visualization
AV = AutoViz_Class()
auto_viz = AV.AutoViz('https://raw.githubusercontent.com/Ivasquez2003/Data-Mining-Project

```

```
Shape of your Data Set loaded: (614, 13)
```

```

#####
##### C L A S S I F Y I N G   V A R I A B L E S #####
#####
Classifying variables in data set...
  Number of Numeric Columns = 3
  Number of Integer-Categorical Columns = 1
  .. .

```

```

Number of String-Categorical Columns = 2
Number of Factor-Categorical Columns = 0
Number of String-Boolean Columns = 5
Number of Numeric-Boolean Columns = 1
Number of Discrete String Columns = 0
Number of NLP String Columns = 0
Number of Date Time Columns = 0
Number of ID Columns = 1
Number of Columns to Delete = 0

```

13 Predictors classified...

1 variable(s) removed since they were ID or low-information variables

List of variables removed: ['Loan\_ID']

To fix these data quality issues in the dataset, import FixDQ from autoviz...

All variables classified into correct types.

	Data Type	Missing Values%	Unique Values%	Minimum Value	Maximum Value	DQ Issue
<b>Loan_ID</b>	object	0.000000	100			Possible ID column: drop before modeling step.
<b>Gender</b>	object	2.117264	0			13 missing values. Impute them with mean, median, mode, or a constant value such as 123., Mixed dtypes: has 2 different data types: object, float,
<b>Married</b>	object	0.488599	0			3 missing values. Impute them with mean, median, mode, or a constant value such as 123., Mixed dtypes: has 2 different data types: object, float,
<b>Dependents</b>	object	2.442997	0			15 missing values. Impute them with mean, median, mode, or a constant value such as 123., Mixed dtypes: has 2 different data types: object, float,
<b>Education</b>	object	0.000000	0			No issue

							32 missing values. Impute them with mean, median, mode, or a constant value such as 123., Mixed dtypes: has 2 different data types: object, float,
<b>Self_Employed</b>	object	5.211726	0				Column has 50 outliers greater than upper bound (10171.25) or lower than lower bound(-1498.75). Cap them or remove them.
<b>ApplicantIncome</b>	int64	0.000000	82	150.000000	81000.000000		Column has 18 outliers greater than upper bound (5743.12) or lower than lower bound(-3445.88). Cap them or remove them.
<b>CoapplicantIncome</b>	float64	0.000000	NA	0.000000	41667.000000		22 missing values. Impute them with mean, median, mode, or a constant value such as 123., Column has 39 outliers greater than upper bound (270.00) or lower than lower bound(-2.00). Cap them or remove them.
<b>LoanAmount</b>	float64	3.583062	NA	9.000000	700.000000		14 missing values. Impute them with mean, median, mode, or a constant value such as 123., Column has 88 outliers greater
<b>Loan_Amount_Term</b>	float64	2.280130	NA	12.000000	480.000000		

than upper  
bound (360.00)  
or lower than  
lower  
bound(360.00).  
Cap them or  
remove them.

50 missing  
values. Impute  
them with mean,  
median, mode, or  
a constant value  
such as 123.

**Credit\_History** float64 8.143322 0

**Property\_Area** object 0.000000 0

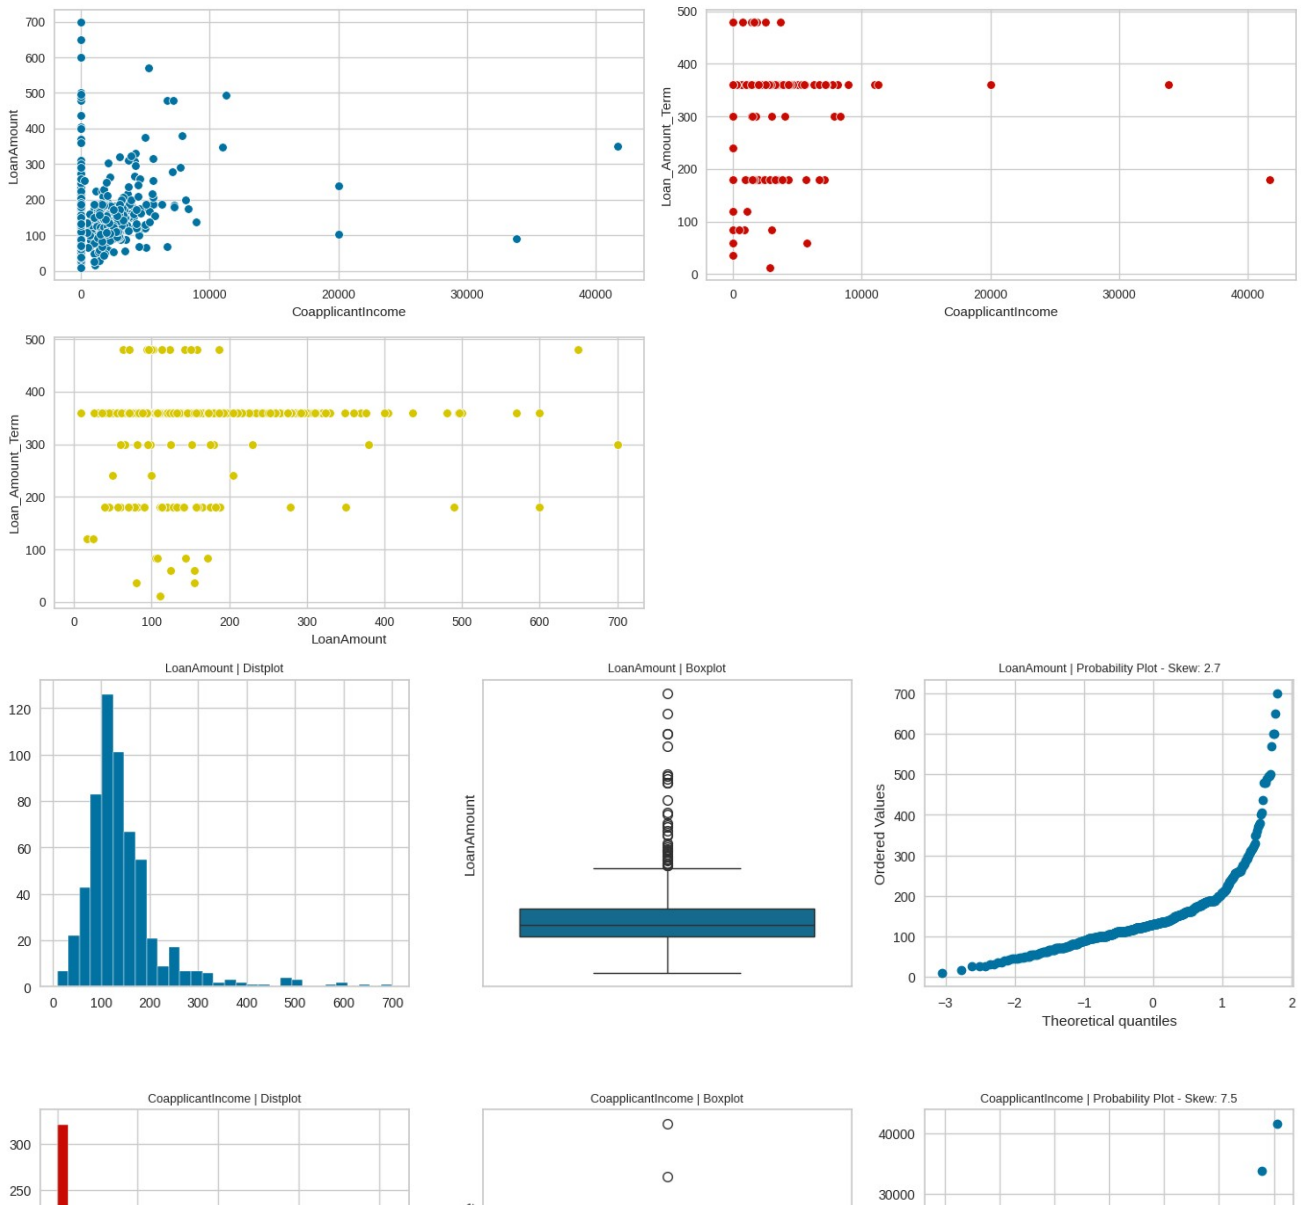
**Loan\_Status** object 0.000000 0

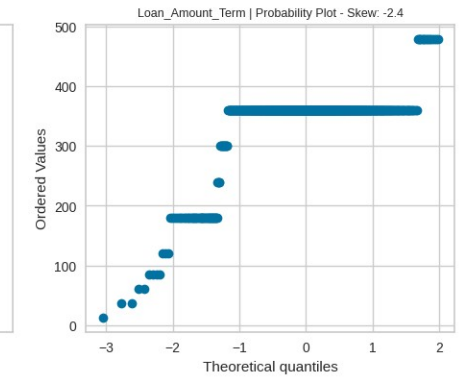
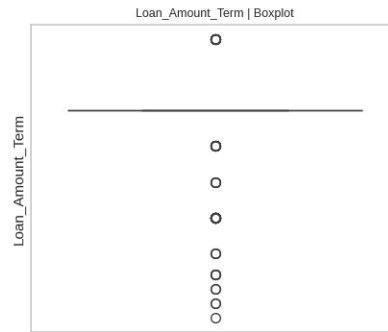
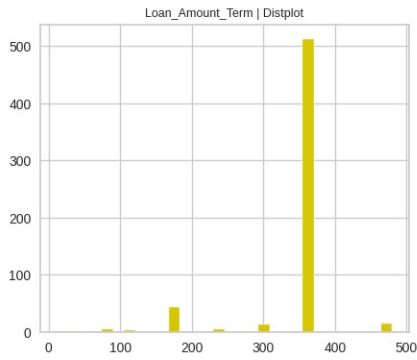



No issue

No issue



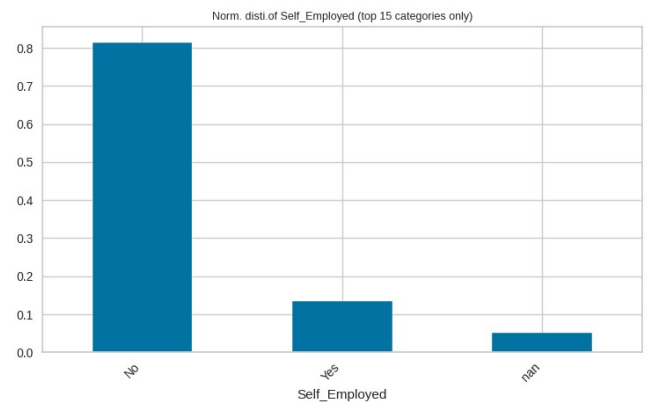
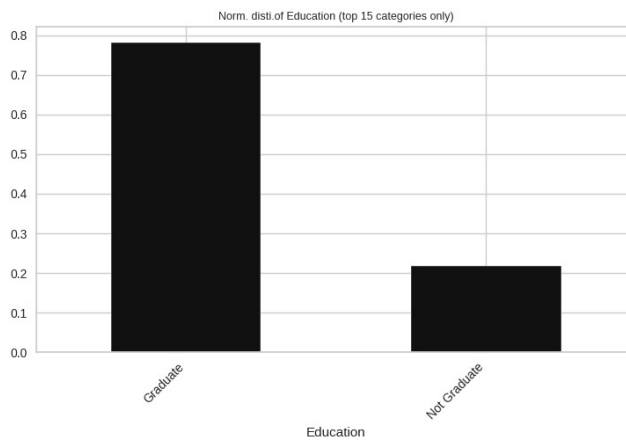
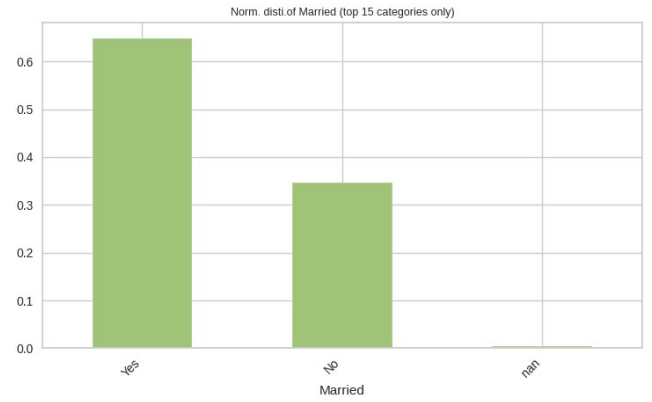
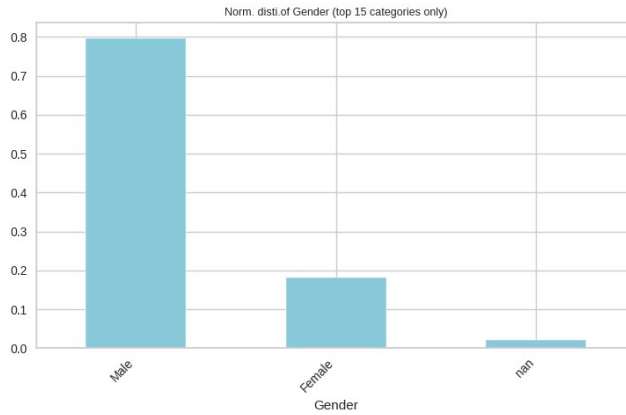
Number of All Scatter Plots = 6

Pair-wise Scatter Plot of all Continuous Variables






Histograms and Normalized distributions of all variables















Violin Plot of all Continuous Variables



















Heatmap of all Numeric Variables including target:











All Plots done

Time to run AutoViz = 19 seconds

##### AUTO VISUALIZATION Completed #####



## ✓ Preparing Data

```
# Drop columns with many missing values or irrelevant
df_clean = df.drop(columns=['Loan_ID'])

# Fill missing values
df_clean['Gender'] = df_clean['Gender'].fillna(df_clean['Gender'].mode()[0])
df_clean['Married'] = df_clean['Married'].fillna(df_clean['Married'].mode()[0])
df_clean['Dependents'] = df_clean['Dependents'].fillna(df_clean['Dependents'].mode()[0])
df_clean['Self_Employed'] = df_clean['Self_Employed'].fillna(df_clean['Self_Employed'].mode()[0])
df_clean['LoanAmount'] = df_clean['LoanAmount'].fillna(df_clean['LoanAmount'].median())
df_clean['Loan_Amount_Term'] = df_clean['Loan_Amount_Term'].fillna(df_clean['Loan_Amount_Term'].mode()[0])
df_clean['Credit_History'] = df_clean['Credit_History'].fillna(df_clean['Credit_History'].mode()[0])

df_clean.dropna(inplace=True)
df_clean.head()
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	Male	No	0	Graduate	No	5849	
1	Male	Yes	1	Graduate	No	4583	
2	Male	Yes	0	Graduate	Yes	3000	
3	Male	Yes	0	Not Graduate	No	2583	
4	Male	No	0	Graduate	No	6000	

## ✓ Model

```
# Setup PyCaret
clf1 = setup(data=df_clean, target='Loan_Status', session_id=123)

# Compare all models
compare_models()

compare_models(include=['rf', 'xgboost'])
```

```
compare_models(include=[ 'rf', 'xgboost' ])
```

	Description	Value
0	Session id	123
1	Target	Loan_Status
2	Target type	Binary
3	Target mapping	N: 0, Y: 1
4	Original data shape	(614, 12)
5	Transformed data shape	(614, 17)
6	Transformed train set shape	(429, 17)
7	Transformed test set shape	(185, 17)
8	Numeric features	5
9	Categorical features	6
10	Preprocess	True
11	Imputation type	simple
12	Numeric imputation	mean
13	Categorical imputation	mode
14	Maximum one-hot encoding	25
15	Encoding method	None
16	Fold Generator	StratifiedKFold
17	Fold Number	10
18	CPU Jobs	-1
19	Use GPU	False
20	Log Experiment	False
21	Experiment Name	clf-default-name
22	USI	8a69

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
<b>rf</b>	Random Forest Classifier	0.7904	0.7597	0.7904	0.7935	0.7717	0.4508	0.4822	1.1320
<b>xgboost</b>	Extreme Gradient Boosting	0.7693	0.7645	0.7693	0.7637	0.7576	0.4208	0.4344	0.8130

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='sqrt',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_samples_leaf=1,
                        min_samples_split=2, n_estimators=100, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

Model Comparison:

Random Forest outperformed XGBoost in overall accuracy (0.7904 vs 0.7683), precision (0.7935 vs 0.7637), and F1-score (0.7717 vs 0.7576). XGBoost achieved a slightly higher AUC score (0.7645 vs 0.7597), indicating it may perform slightly better in distinguishing between classes under imbalance.

Final Decision: Random Forest was selected as the final model due to its superior performance in accuracy, precision, and F1-score — key metrics for the loan approval prediction.

Interpretation: The most influential features in the Random Forest model were:

- **Credit\_History:** Applicants with good credit history had significantly higher approval chances.
- **LoanAmount:** Higher requested amounts generally correlated with a lower approval probability.
- **ApplicantIncome:** Income had moderate influence, especially when considered with Coapplicant Income.

## ✓ Evaluate

```
# Create and evaluate a Random Forest model
rf_model = create_model('rf')
evaluate_model(rf_model)
predict_model(rf_model)
```

	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
Fold							
0	0.8140	0.8077	0.8140	0.8212	0.7945	0.4926	0.5327
1	0.6977	0.6859	0.6977	0.6694	0.6733	0.1957	0.2058
2	0.7674	0.6436	0.7674	0.7743	0.7294	0.3323	0.3931
3	0.7907	0.8449	0.7907	0.8390	0.7489	0.3828	0.4865
4	0.7907	0.8154	0.7907	0.7867	0.7882	0.4928	0.4936
5	0.7907	0.7241	0.7907	0.7888	0.7763	0.4749	0.4965
6	0.7907	0.7303	0.7907	0.7888	0.7763	0.4749	0.4965

<b>7</b>	0.8372	0.7734	0.8372	0.8448	0.8260	0.5916	0.6185
<b>8</b>	0.7674	0.6749	0.7674	0.7645	0.7466	0.4044	0.4330
<b>9</b>	0.8571	0.8966	0.8571	0.8571	0.8571	0.6658	0.6658
<b>Mean</b>	0.7904	0.7597	0.7904	0.7935	0.7717	0.4508	0.4822
<b>Std</b>	0.0411	0.0773	0.0411	0.0512	0.0487	0.1251	0.1192

Plot Type:

Pipeline Plot	Hyperparameters	AUC	Confusion Matrix
Threshold	Precision Recall	Prediction Error	Class Report
Feature Selection	Learning Curve	Manifold Learning	Calibration Curve
Validation Curve	Dimensions	Feature Importance	Feature Importance ...
Decision Boundary	Lift Chart	Gain Chart	Decision Tree
KS Statistic Plot			



	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
<b>0</b>	Random Forest Classifier	0.7892	0.7748	0.7892	0.8015	0.7622	0.4271	0.4797
	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic	
<b>354</b>	Female	Yes	0	Graduate	No	2423		
<b>426</b>	Female	No	1	Not Graduate	No	4606		
<b>611</b>	Male	Yes	1	Graduate	No	8072		
<b>214</b>	Male	Yes	0	Graduate	No	3173		
<b>72</b>	Male	No	0	Graduate	No	3500		
...	...	...	...	...	...	...		
<b>574</b>	Male	Yes	3+	Graduate	No	6406		
<b>42</b>	Male	Yes	0	Graduate	No	2400		
<b>50</b>	Female	Yes	0	Not Graduate	No	1928		
<b>189</b>	Male	Yes	0	Graduate	No	9328		
<b>215</b>	Male	Yes	3+	Not Graduate	No	3850		

185 rows × 14 columns



